A brief tour of R/qtl

Karl W Broman

Department of Biostatistics and Medical Informatics
University of Wisconsin – Madison

https://rqtl.org

21 March 2012

Overview of R/qtl

R/qtl is an extensible, interactive environment for mapping quantitative trait loci (QTL) in experimental crosses. It is implemented as an add-on package for the freely available and widely used statistical language/software R (see www.r-project.org). The development of this software as an add-on to R allows us to take advantage of the basic mathematical and statistical functions, and powerful graphics capabilities, that are provided with R. Further, the user will benefit by the seamless integration of the QTL mapping software into a general statistical analysis program. Our goal is to make complex QTL mapping methods widely accessible and allow users to focus on modeling rather than computing.

A key component of computational methods for QTL mapping is the hidden Markov model (HMM) technology for dealing with missing genotype data. We have implemented the main HMM algorithms, with allowance for the presence of genotyping errors, for backcrosses, intercrosses, and phase-known four-way crosses.

The current version of R/qtl includes facilities for estimating genetic maps, identifying genotyping errors, and performing single-QTL genome scans and two-QTL, two-dimensional genome scans, by interval mapping (with the EM algorithm), Haley-Knott regression, and multiple imputation. All of this may be done in the presence of covariates (such as sex, age or treatment). One may also fit higher-order QTL models by multiple imputation and Haley-Knott regression.

R/qtl is distributed as source code for Unix or compiled code for Windows or Mac OS X. R/qtl is released under the GNU General Public License, version 3. To download the software, you must agree to the terms in that license.

Overview of R

R is an open-source implementation of the S language. As described on the R-project homepage (www.r-project.org):

> R is a system for statistical computation and graphics. It consists of a language plus a run-time environment with graphics, a debugger, access to certain system functions, and the ability to run programs stored in script files.

> The core of R is an interpreted computer language which allows branching and looping as well as modular programming using functions. Most of the user-visible functions in R are written in R. It is possible for the user to interface to procedures written in the C, C++, or FORTRAN languages for efficiency. The R distribution contains functionality for a large number of statistical procedures. Among these are: linear and generalized linear models, nonlinear regression models, time series analysis, classical parametric and nonparametric tests, clustering and smoothing. There is also a large set of functions which provide a flexible graphical environment for creating various kinds of data presentations. Additional modules are available for a variety of specific purposes.

R is freely available for Windows, Unix and Mac OS X, and may be downloaded from the Comprehensive R Archive Network (CRAN; cran.r-project.org).

Learning R may require a formidable investment of time, but it will definitely be worth the effort. Numerous free documents on getting started with R are available on CRAN. In additional, several books are available. The most important book on R is Venables and Ripley (2002) *Modern Applied Statistics with S*, 4th edition. Dalgaard (2002) *Introductory Statistics with R* provides a more gentle introduction.

Citation for R/qtl

To cite R/qtl in publications, use

Broman KW, Wu H, Sen S, Churchill GA (2003) R/qtl: QTL mapping in experimental crosses. Bioinformatics 19:889-890

Selected R/qtl functions

Sample data	badorder	An intercross with misplaced markers
	bristle3	Data on bristle number for Drosophila chromosome 3
	bristleX	Data on bristle number for Drosophila X chromosome
	fake.4way	Simulated data for a 4-way cross
	fake.bc	Simulated data for a backcross
	fake.f2	Simulated data for an F_2 intercross
	hyper	Backcross data on salt-induced hypertension
	listeria	Intercross data on Listeria monocytogenes susceptibility
	map10	A genetic map modeled after the mouse genome (10 cM spacing)
Input/output	read.cross	Read data for a QTL experiment
	write.cross	Write data for a QTL experiment to a file
Simulation	sim.cross	Simulate a QTL experiment
	sim.map	Generate a genetic map
Summaries	qtlversion	Gives the version number of installed R/qtl package
	plot.cross	Plot various features of a cross object
	plotMissing	Plot grid of missing genotypes
	geno.image	Plot grid with colored pixels representing different genotypes
	plotPheno	Histogram or bar plot of a phenotype
	plotInfo	Plot the proportion of missing genotype data
	summary.cross	Print summary of QTL experiment
	summaryMap	Print summary of a genetic map
	nchr, nind, nmar, nphe, totmar, nmissing, ntyped	
	find.pheno	Find the column number for a particular phenotype
	find.marker	Find the marker closest to a specified position
	find.flanking	Find the markers flanking a particular position
	find.pseudomarker	Find the name of the grid position closest to a particular position
	find.markerpos	Find the map positions of a marker
Data manipulation	clean.cross	Remove intermediate calculations from a cross
	drop.markers	Remove a list of markers
	drop.nullmarkers	Remove markers without data
	fill.geno	Fill in holes in genotype data by imputation or Viterbi
	strip.partials	Replace partially informative genotypes with missing values
	pull.map	Pull out the genetic map from a cross
	pull.geno	Pull out the genotype data as a matrix
	pull.pheno	Pull out a phenotype
	replace.map	Replace the genetic map of a cross
	jittermap	Jitter marker positions slightly so that no two coincide
	subset.cross	Select a subset of chromosomes and/or individuals from a cross
	c.cross	Combine two crosses into one object
	switch.order	Switch the order of markers on a chromosome
	movemarker	Move a marker from one chromosome to another
HMM engine	argmax.geno	Reconstruct underlying genotypes by the Viterbi algorithm
	calc.genoprob	Calculate conditional genotype probabilities
	sim.geno	Simulate genotypes given observed marker data
Diagnostics	geno.table	Create table of genotype distributions
	geno.crosstab	Create cross-tabulation of genotypes at two markers
	checkAlleles	Identify markers with potentially switched alleles
	calc.errorlod	Calculate Lincoln & Lander (1992) error LOD scores
	top.errorlod	List genotypes with highest error LOD values
	plotGeno	Plot observed genotypes, flagging likely errors
	comparecrosses	Compare two cross objects, to see if they are the same
	comparegeno	Calculate proportion of matching genotypes for each pair of individuals

Selected R/qtl functions (continued)

Genetic mapping	est.rf	Estimate pairwise recombination fractions
	plotRF	Plot recombination fractions
	est.map	Estimate genetic map
	plotMap	Plot genetic map(s)
	summaryMap	Print summary of a genetic map
	ripple	Assess marker order by permuting groups of adjacent markers
	summary.ripple	Print summary of ripple output
	compareorder	Compare two orderings of markers on a chromosome
	tryallpositions	Test all possible positions for a marker
QTL mapping	scanone	Genome scan with a single QTL model
	scantwo	Two-dimensional genome scan with a two-QTL model
	lodint	Calculate a LOD support interval
	bayesint	Calculate an approximate Bayes credible interval
	scanoneboot	Non-parametric bootstrap to obtain a confidence interval for QTL location
	plot.scanone	Plot output for a one-dimensional genome scan
	add.threshold	Add a horizontal line at a LOD threshold to a genome scan plot
	plot.scantwo	Plot output for a two-dimensional genome scan
	summary.scanone	Print summary of scanone output
	summary.scantwo	Print summary of scantwo output
	max.scanone	Maximum peak in scanone output
	max.scantwo	Maximum peak in scantwo output
	effectplot	Plot phenotype means of genotype groups defined by 1 or 2 markers
	effectscan	Plot estimated QTL effects across the whole genome
	plotPXG	Like effectplot, but as a dot plot of the phenotypes
Multiple QTL models	makeqtl	Make a qtl object for use by fitqtl
	fitqtl	Fit a multiple QTL model
	summary.fitqtl	Get summary of the result of fitqtl
	scanqtl	Perform a multi-dimensional genome scan
	refineqtl	Refine the QTL locations in a multiple QTL model
	plotLodProfile	Plot 1-dimensional LOD profiles for a multiple QTL model
	addqtl	Scan for an additional QTL, in a multiple-QTL model
	addpair	Scan for an additional pair of QTL, in a multiple-QTL model
	addint	Add pairwise interactions, one at a time, in a multiple-QTL model
	summary.qtl	Print a summary of a QTL object
	plot.qtl	Plot the QTL locations on the genetic map
	addtoqtl	Add to a QTL object
	dropfromqtl	Drop a QTL from a QTL object
	replaceqtl	Replace a QTL location in a QTL object with a different position
	reorderqtl	Reorder the QTL in a QTL object
	cim	A (relatively crude) implementation of Composite Interval Mapping
	stepwiseqtl	Stepwise selection for multiple QTL
	calc.penalties	Calculate penalties for use with stepwiseqtl
	plotModel	Plot a graphical representation of a multiple-QTL model

Preliminaries

Use of the R/qtl package requires considerable knowledge of the R language/environment. We hope that the examples presented here will be understandable with little prior knowledge of R, especially because we neglect to explain the syntax of R. Several books, as well as some free documents, are available to assist the user in learning R; see the R project website cited above. We assume here that the user is running either Windows or Mac OS X.

1. To start R, double-click its icon.

2. To exit, type:

   ```
   q()
   ```

   Click yes or no to save or discard your work.

3. R keeps all of your work in RAM. If R should crash, all will be lost, and you will have to start from the beginning. The function `save.image` can be used to save your work to disk as you go along, so that, should R crash, you won't have to start from scratch. You would type:

   ```
   save.image()
   ```

4. Load the R/qtl package:

   ```
   library(qtl)
   ```

5. View the objects in your workspace:

   ```
   ls()
   ```

6. The best way to get help on the functions and data sets in R (and in R/qtl) is via the html version of the help files. One way to get access to this is to type

   ```
   help.start()
   ```

   This should open a browser with the main help menu. If you then click on Packages → qtl, you can see all of the available functions and datasets in R/qtl. For example, look at the help file for the function `read.cross`.

   An alternative method to view this help file is to type one of the following:

   ```
   help(read.cross)
   ?read.cross
   ```

   The html version of the help files are somewhat easier to read, and allow use of hotlinks between different functions.

7. All of the code in this tutorial is available as a file from which you may copy and paste into R, if you prefer that to typing. Type the following within R to get access to the file:

   ```
   url.show("https://rqtl.org/rqtltour.R")
   ```

Data import

A difficult first step in the use of most data analysis software is the import of data. With R/qtl, one may import data in several different formats by use of the function `read.cross`. (Example data files are available at https://rqtl.org/sampledata.) The internal data structure used by R/qtl is rather complicated, and is described in the help file for `read.cross`. (Also see Example 6 on page 19.) We won't discuss data import any further here, except to say that the comma-delimited format (`"csv"`) is recommended. If you have trouble importing data, send an email to Karl Broman (`broman@wisc.edu`), attaching examples of your data files. (Such data will be kept confidential.)

Example 1: Hypertension

As a first example, we consider data from an experiment on hypertension in the mouse (Sugiyama et al., Genomics 71:70-77, 2001), kindly provided by Bev Paigen and Gary Churchill.

1. First, get access to the data, see that it is in your workspace, and view its help file. These data are included with the R/qtl package, and so you can get access to the data with the function `data()`. (Remember that you first need to load the R/qtl package via `library(qtl)`.)

   ```
   data(hyper)
   ls()
   ?hyper
   ```

2. We will postpone discussion of the internal data structure used by R/qtl until later. For now we'll just say that the data `hyper` has "class" `"cross"`. The function `summary.cross` prints summary information on such data. We can call that function directly, or we may simply use `summary` and the data is sent to the appropriate function according to its class.

```
summary(hyper)
```

Several other utility functions are available for getting summary information on the data. Hopefully these are self-explanatory.

```
nind(hyper)
nphe(hyper)
nchr(hyper)
totmar(hyper)
nmar(hyper)
```

3. Plot a summary of these data.

```
plot(hyper)
```

In the upper left, black pixels indicate missing genotype data. Note that one marker has no genotype data. In the upper right, the genetic map of the markers is shown. In the lower left, a histogram of the phenotype is shown.

The Windows version of R has a slick method for recording graphs, so that one may page up and down through a series of plots. To initiate this, click (on the menu bar) History → Recording.

We may plot the individual components of the above multi-plot figure as follows.

```
plotMissing(hyper)
plotMap(hyper)
plotPheno(hyper, pheno.col=1)
```

We can plot the genetic map with marker names, but they can be rather difficult to read. The following code plots the map with marker names for chr 1, 4, 6, 7 and 15.

```
plotMap(hyper, chr=c(1, 4, 6, 7, 15), show.marker.names=TRUE)
```

4. Note the odd pattern of missing data; we may make this missing data plot with the individuals ordered according to the value of their phenotype.

```
plotMissing(hyper, reorder=TRUE)
```

We see that, for most markers, only individuals with extreme phenotypes were genotyped. At many markers (in regions of interest), markers were typed only on recombinant individuals.

5. The function `drop.nullmarkers` may be used to remove markers that have no genotype data (such as the marker on chr 14). A call to `totmar` will show that there are now 173 markers (rather than 174, as there were initially).

```
hyper <- drop.nullmarkers(hyper)
totmar(hyper)
```

6. Estimate recombination fractions between all pairs of markers, and plot them. This also calculates LOD scores for the test of $H_0: r = 1/2$. The plot of the recombination fractions can be either with recombination fractions in the upper part and LOD scores below, or with just recombination fractions or just LOD scores. Note that red corresponds to a small recombination fraction or a big LOD score, while blue is the reverse. Gray indicates missing values.

```
hyper <- est.rf(hyper)
plotRF(hyper)
plotRF(hyper, chr=c(1,4))
```

There are some very strange patterns in the recombination fractions, but this is due to the fact that some markers were typed largely on recombinant individuals.

For example, on chr 6, the tenth marker shows a high recombination fraction with all other markers on the chromosome, but a plot of the missing data shows that this marker was typed only on a selected number of individuals (largely those showing recombination events across the interval).

```
plotRF(hyper, chr=6)
plotMissing(hyper, chr=6)
```

7. Re-estimate the genetic map (keeping the order of markers fixed), and plot the original map against the newly estimated one.

```
newmap <- est.map(hyper, error.prob=0.01)
plotMap(hyper, newmap)
```

We see some map expansion, especially on chr 6, 13 and 18. It is questionable whether we should replace the map or not. Keep in mind that the previous map locations are based on a limited number of meioses. If one wished to replace the genetic map with the estimated one, it could be done as follows:

```
hyper <- replace.map(hyper, newmap)
```

This replaces the map in the `hyper` data with `newmap`.

8. We now turn to the identification of genotyping errors. In the following, we calculate the error LOD scores of Lincoln and Lander (1992). A LOD score is calculated for each individual at each marker; large scores indicate likely genotyping errors.

```
hyper <- calc.errorlod(hyper, error.prob=0.01)
```

This calculates the genotype error LOD scores and inserts them into the `hyper` object.

The function `top.errorlod` gives a list of genotypes that may be in error. Error LOD scores < 4 can probably be ignored.

```
top.errorlod(hyper)
```

Note that the results will be different, depending on whether you used `replace.map` above. If you did, you will get an indication of potential errors on chr 16 (and a few on chr 13). If you didn't, you will get a very long list of potential errors on chr 1, 11, 15, 16 and 17.

9. The function `plotGeno` may be used to inspect the observed genotypes for a chromosome, with likely genotyping errors flagged. Of course, it's difficult to look at too many individuals at once. Note that white = AA and black = AB (for a backcross).

```
plotGeno(hyper, chr=16, ind=c(24:34, 71:81))
```

We don't have any utilities for fixing any apparent errors; it would be best to go back to the raw data. (Of course, you should edit a copy of the file; never discard the primary data.)

10. The function `plotInfo` plots a measure of the proportion of missing genotype information in the genotype data. The missing information is calculated in two ways: as entropy, or via the variance of the conditional genotypes, given the observed marker data. (See the help file, using `?plotInfo`.)

```
plotInfo(hyper)
plotInfo(hyper, chr=c(1,4,15))
plotInfo(hyper, chr=c(1,4,15), method="entropy")
plotInfo(hyper, chr=c(1,4,15), method="variance")
```

11. We now, finally, get to QTL mapping.

The core of R/qtl is a set of functions which make use of the hidden Markov model (HMM) technology to calculate QTL genotype probabilities, to simulate from the joint genotype distribution and to calculate the most likely sequence of underlying genotypes (all conditional on the observed marker data). This is done in a quite general way, with possible allowance for the presence of genotyping errors. Of course, for convenience we assume no crossover interference.

The function `calc.genoprob` calculates QTL genotype probabilities, conditional on the available marker data. These are needed for most of the QTL mapping functions. The argument `step` indicates the step size (in cM) at which the probabilities are calculated, and determines the step size at which later LOD scores are calculated.

```
hyper <- calc.genoprob(hyper, step=1, error.prob=0.01)
```

We may now use the function `scanone` to perform a single-QTL genome scan with a normal model. We may use maximum likelihood via the EM algorithm (Lander and Botstein 1989) or use Haley-Knott regression (Haley and Knott 1992).

```
out.em <- scanone(hyper)
out.hk <- scanone(hyper, method="hk")
```

We may also use the multiple imputation method of Sen and Churchill (2001). This requires that we first use `sim.geno` to simulate from the joint genotype distribution, given the observed marker data. Again, the argument `step` indicates

the step size at which the imputations are performed and determines the step size at which LOD scores will be calculated. The `n.draws` indicates the number of imputations to perform. Larger values give more precise results but require considerably more computer memory and computation time.

```
hyper <- sim.geno(hyper, step=2, n.draws=16, error.prob=0.01)
out.imp <- scanone(hyper, method="imp")
```

12. The output of scanone has class `"scanone"`; the function `summary.scanone` displays the maximum LOD score on each chromosome for which the LOD exceeds a specified threshold.

```
summary(out.em)
summary(out.em, threshold=3)
summary(out.hk, threshold=3)
summary(out.imp, threshold=3)
```

13. The function `max.scanone` returns just the highest peak from output of `scanone`.

```
max(out.em)
max(out.hk)
max(out.imp)
```

14. We may also plot the results. `plot.scanone` can plot up to three genome scans at once, provided that they conform appropriately. Alternatively, one may use the argument `add`.

```
plot(out.em, chr=c(1,4,15))
plot(out.em, out.hk, out.imp, chr=c(1,4,15))
plot(out.em, chr=c(1,4,15))
plot(out.hk, chr=c(1,4,15), col="blue", add=TRUE)
plot(out.imp, chr=c(1,4,15), col="red", add=TRUE)
```

15. The function `scanone` may also be used to perform a permutation test to get a genome-wide LOD significance threshold. For Haley-Knott regression, this can be quite fast.

```
operm.hk <- scanone(hyper, method="hk", n.perm=1000)
```

The permutation output has class `"scanoneperm"`. The function `summary.scanoneperm` can be used to get significance thresholds.

```
summary(operm.hk, alpha=0.05)
```

In addition, if the permutations results are included in a call to `summary.scanone`, you can estimated genome-scan-adjusted p-values for inferred QTL, and can get a report of all chromosomes meeting a certain significance level, with the corresponding LOD threshold calculated automatically.

```
summary(out.hk, perms=operm.hk, alpha=0.05, pvalues=TRUE)
```

16. We should mention at this point that the function `save.image` may be used to save your workspace to disk. If R crashes, you will wish you had used this.

```
save.image()
```

17. The function `scantwo` performs a two-dimensional genome scan with a two-QTL model. For every pair of positions, it calculates a LOD score for the full model (two QTL plus interaction) and a LOD score for the additive model (two QTL but no interaction). This be quite time consuming, and so you may wish to do the calculations on a coarser grid.

```
hyper <- calc.genoprob(hyper, step=5, error.prob=0.01)
out2.hk <- scantwo(hyper, method="hk")
```

One can also use `method="em"` or `method="imp"`, but they are even more time consuming.

18. The output of scantwo has class `"scantwo"`; there are functions for obtaining summaries and plots, of course.

The summary function considers each pair of chromosomes, and calculates the maximum LOD score for the full model (M_f) and the maximum LOD score for the additive model (M_a). These two models are allowed to be maximized at different positions. We futher calculate a LOD score for a test of epistasis, $M_i = M_f - M_a$, and two LOD scores that concern evidence for a second QTL: M_{fv1} is the LOD score comparing the full model to the best single-QTL model and M_{av1} is the LOD score comparing the additive model to the best single-QTL model.

In the summary, we must provide five thresholds, for M_f, M_{fv1}, M_i, M_a, and M_{av1}, respectively. Call these T_f, T_{fv1}, T_i, T_a, and T_{av1}. We then report those pairs of chromosomes for which at least one of the following holds:

- $M_f \geq T_f$ and ($M_{fv1} \geq T_{fv1}$ or $M_i \geq T_i$)
- $M_a \geq T_a$ and $M_{av1} \geq T_{av1}$

The thresholds can be obtained by a permutation test (see below), but this is extremely time-consuming. For a mouse backcross, we suggest the thresholds (6.0, 4.7, 4.4, 4.7, 2.6) for the full, conditional-interactive, interaction, additive, and conditional-additive LOD scores, respectively. For a mouse intercross, we suggest the thresholds (9.1, 7.1, 6.3, 6.3, 3.3) for the full, conditional-interactive, interaction, additive, and conditional-additive LOD scores, respectively. These were obtained by 10,000 simulations of crosses with 250 individuals, markers at a 10 cM spacing, and analysis by Haley-Knott regression.

```
summary(out2.hk, thresholds=c(6.0, 4.7, 4.4, 4.7, 2.6))
```

The appropriate decision rule is not yet completely clear. I am inclined to ignore M_i and to choose genome-wide thresholds for the other four based on a permutation, using a common significance level for all four. M_i would be ignored if we gave it a very large threshold, as follows.

```
summary(out2.hk, thresholds=c(6.0, 4.7, Inf, 4.7, 2.6))
```

19. Plots of `scantwo` results are created via `plot.scantwo`.

```
plot(out2.hk)
plot(out2.hk, chr=c(1,4,6,15))
```

By default, the upper-left triangle contains epistasis LOD scores and the lower-right triangle contains the LOD scores for the full model. The color scale on the right indicates separate scales for the epistasis and joint LOD scores (on the left and right, respectively).

20. The function `max.scantwo` returns the two-locus positions with the maximum LOD score for the full and additive models.

```
max(out2.hk)
```

21. One may also use `scantwo` to perform permutation tests in order to obtain genome-wide LOD significance thresholds. These can be extremely time consuming, though with the Haley-Knott regression and multiple imputation methods, there is a trick that may be used in some cases to dramatically speed things up. So we'll try 100 permutations by the Haley-Knott regression method and hope that your computer is sufficiently fast.

```
operm2.hk <- scantwo(hyper, method="hk", n.perm=100)
```

We can again use `summary` to get LOD thresholds.

```
summary(operm2.hk)
```

And again these may be used in the summary of the `scantwo` output to calculate thresholds and p-values. If you want to ignore the LOD score for the interaction in the rule about what chromosome pairs to report, give $\alpha = 0$, corresponding to a threshold $T = \infty$.

```
summary(out2.hk, perms=operm2.hk, pvalues=TRUE,
        alphas=c(0.05, 0.05, 0, 0.05, 0.05))
```

You can't really trust these results. Haley-Knott regression performs poorly in the case of selective genotyping (as with the `hyper` data). Standard interval mapping or imputation would be better, but Haley-Knott regression has the advantage of speed, which is the reason we use it here.

22. Finally, we consider the fit of multiple-QTL models. Currently, only multiple imputation and Haley-Knott regression has been implemented. We use multiple imputation here, as Haley-Knott regression performs poorly in the case of selective genotyping, which was used for the `hyper` data. We first create a QTL object using the function `makeqtl`, with five QTL at specified, fixed positions.

```
chr <- c(1, 1, 4, 6, 15)
pos <- c(50, 76, 30, 70, 20)
qtl <- makeqtl(hyper, chr, pos)
```

Finally, we use the function `fitqtl` to fit a model with five QTL, and allowing the QTL on chr 6 and 15 to interact.

```
my.formula <- y ~ Q1 + Q2 + Q3 + Q4 + Q5 + Q4:Q5
out.fitqtl <- fitqtl(hyper, qtl=qtl, formula=my.formula)
summary(out.fitqtl)
```

See Example 5 (page 14) for a thorough discussion of the multiple QTL mapping methods in R/qtl.

23. You may wish to clean up your workspace before we move on to the next example.

```
ls()
rm(list=ls())
```

Example 2: Genetic mapping

R/qtl includes some utilities for estimating genetics maps and checking marker orders. In this example, we describe the use of these utilities.

1. Get access to some sample data. This is simulated data with some errors in marker order.

```
data(badorder)
summary(badorder)
plot(badorder)
```

2. Estimate recombination fractions between all pairs of markers, and plot them.

```
badorder <- est.rf(badorder)
plotRF(badorder)
```

It appears that markers on chr 2 and 3 have been switched.

Also note that, if we look more closely at the recombination fractions for chr 1, there seem to be some errors in marker order.

```
plotRF(badorder, chr=1)
```

3. Re-estimate the genetic map.

```
newmap <- est.map(badorder, verbose=TRUE)
plotMap(badorder, newmap)
```

This really shows the problems on chr 2 and 3.

4. Fix the problems on chr 2 and 3. First, we look more closely at the recombination fractions for these chromosoems

```
plotRF(badorder, chr=2:3)
```

We need to move the sixth marker on chr 2 to chr 3, and the fifth marker on chr 3 to chr 2. We need to figure out which markers these are.

```
pull.map(badorder, chr=2)
pull.map(badorder, chr=3)
```

Now we can use the function `movemarker` to move the markers. It seems like they should be exactly switched.

```
badorder <- movemarker(badorder, "D2M937", 3, 48)
badorder <- movemarker(badorder, "D3M160", 2, 28.8)
```

Now look at the recombination fractions again.

```
plotRF(badorder, chr=2:3)
```

5. We can check the marker order on chr 1. The function `ripple` will consider all permutations of a sliding window of adjacent markers. A quick-and-dirty approach is to count the number of obligate crossovers for each possible order, to find the order with the minimum number of crossovers. A more refined, but also more computationally intensive, approach is to re-estimate the genetic map for each order, calculating LOD scores ($\log_{10}$ likelihood ratios) relative to the initial order. (This may be done with allowance for the presence of genotyping errors.) The default approach is the quick-and-dirty method.

The following checks the marker order on chr 1, permuting groups of six contiguous markers.

```
rip1 <- ripple(badorder, chr=1, window=6)
summary(rip1)
```

In the summary output, markers 9–11 clearly need to be flipped. There also seems to be a problem with the order of markers 4–6.

6. The following performs the likelihood analysis, permuting groups of three adjacent markers, assuming a genotyping error rate of 1%. It's considerably slower, but more trustworthy.

```
rip2 <- ripple(badorder, chr=1, window=3, err=0.01, method="likelihood")
summary(rip2)
```

Note that positive LOD scores indicate that the alternate order has a higher likelihood than the original.

7. We can switch the order of markers 9–11 with the function `switch.order` (which works only for a single chromosome) and then re-assess the order. Note that the second row of `rip1` corresponds to the improved order.

```
badorder.rev <- switch.order(badorder, 1, rip1[2,])
rip1r <- ripple(badorder.rev, chr=1, window=6)
summary(rip1r)
```

It looks like the marker pairs (5,6) and (1,2) should each be inverted. We use `switch.order` again, and then check marker order using the likelihood method.

```
badorder.rev <- switch.order(badorder.rev, 1, rip1r[2,])
rip2r <- ripple(badorder.rev, chr=1, window=3, err=0.01)
summary(rip2r)
```

It's probably best to start out using the quick-and-dirty method, with a large window size, to find the marker order with the minimum number of obligate crossovers, and then refine that order using the slower, but more trustworthy, likelihood method.

8. We can look again at the recombination fractions for this chromosome.

```
badorder.rev <- est.rf(badorder.rev)
plotRF(badorder.rev, 1)
```

Example 3: Listeria susceptibility

In order to demonstrate further uses of the function `scanone`, we consider some data on susceptibility to *Listeria monocytogenes* in mice (Boyartchuk et al., Nature Genetics 27:259-260, 2001). These data were kindly provided by Victor Boyartchuk and Bill Dietrich.

1. Get access to the data and view some summaries.

```
data(listeria)
summary(listeria)
plot(listeria)
plotMissing(listeria)
```

Note that in the missing data plot, gray pixels are partially missing genotypes (e.g., a genotype may be known to be either AA or AB, but not which).

The phenotype here is the survival time of a mouse (in hours) following infection with *Listeria monocytogenes*. Individuals with a survival time of 264 hours are those that recovered from the infection.

2. We'll use the log survival time, rather than survival time, so we first need to create a new phenotype, which will end up as the third phenotype (after `sex`).

```
listeria$pheno$logSurv <- log(listeria$pheno[,1])
plot(listeria)
```

3. Estimate pairwise recombination fractions.

```
listeria <- est.rf(listeria)
plotRF(listeria)
plotRF(listeria, chr=c(5,13))
```

4. Re-estimate the genetic map.

```
newmap <- est.map(listeria, error.prob=0.01)
plotMap(listeria, newmap)
listeria <- replace.map(listeria, newmap)
```

5. Investigate genotyping errors; nothing gets flagged with a cutoff of 4, but one genotype is indicated with error LOD ~3.8.

```
listeria <- calc.errorlod(listeria, error.prob=0.01)
top.errorlod(listeria)
top.errorlod(listeria, cutoff=3.5)
plotGeno(listeria, chr=13, ind=61:70, cutoff=3.5)
```

Note that in the plot given by `plotGeno`, for an intercross, white = AA, gray = AB, black = BB, green = AA or AB, and orange = AB or BB.

6. Now on to the QTL mapping. Recall that the phenotype distribution shows a clear departure from the standard assumptions for interval mapping; 30% of the mice survived longer than 264 hours, and were considered recovered from the infection.

One approach for these data is to use the two-part model considered by Boyartchuk et al. (2001). In this model, a mouse with genotype g has probability p_g of surviving the infection. If it does die, its log survival time is assumed to be distributed normal(μ_g, σ^2). Analysis proceeds by maximum likelihood via an EM algorithm. Three LOD scores are calculated. LOD(p, μ) is for the test of the null hypothesis $p_g \equiv p$ and $\mu_g \equiv \mu$. LOD(p) is for the test of the hypothesis $p_g \equiv p$ but the μ are allowed to vary. LOD(μ) is for the test of the hypothesis $\mu_g \equiv \mu$ but the p are allowed to vary.

The function `scanone` will fit the above model when the argument `model="2part"`. One must also specify the argument `upper`, which indicates whether the spike in the phenotype is the maximum phenotype (as it is with this phenotype; take `upper=TRUE`) or the minimum phenotype (take `upper=FALSE`). For this model, only the EM algorithm has been implemented so far.

```
listeria <- calc.genoprob(listeria, step=2)
out.2p <- scanone(listeria, pheno.col=3, model="2part", upper=TRUE)
```

Note the use of the argument `pheno.col` to indicate the phenotype column to use for the analysis. We can also refer to the phenotype column by name: `pheno.col="logSurv"`.

Because the two-part model has three extra parameters, the appropriate LOD threshold is higher—around 4.5 rather than 3.5. The three different LOD curves are in columns 3–5 of the output.

```
summary(out.2p)
summary(out.2p, threshold=4.5)
```

Alternatively, we may use `format="allpeaks"`, in which case it displays the maximum LOD score or each column, with the position at which each was maximized. You may provide either one threshold, which would be applied to all LOD score columns, or a separate threshold for each column.

```
summary(out.2p, format="allpeaks", threshold=3)
summary(out.2p, format="allpeaks", threshold=c(4.5,3,3))
```

7. By default, `plot.scanone` will plot the first LOD score column. Alternatively, we may indicate another column to plot with the `lodcolumn` argument. Or we can plot up to three LOD scores at once by giving a vector.

```
plot(out.2p)
plot(out.2p, lodcolumn=2)
plot(out.2p, lodcolumn=1:3, chr=c(1,5,13,15))
```

Note that the locus on chr 1 shows effect mostly on the mean time-to-death, conditional on death; the locus on chr 5 shows effect mostly on the probability of survival; and the loci on chr 13 and 15 shows some effect on each.

8. Permutation tests may be performed as before. The output will have three columns, corresponding to the three LOD scores.

```
operm.2p <- scanone(listeria, model="2part", pheno.col=3,
                    upper=TRUE, n.perm=25)
summary(operm.2p, alpha=0.05)
```

We may again use the permutation results in `summary.scanone` to have thresholds calculated automatically and to obtain genome-scan-adjusted p-values, but of course we would want to have performed more than 25 permutations.

```
summary(out.2p, format="allpeaks", perms=operm.2p,
        alpha=0.05, pvalues=TRUE)
```

11

9. Alternatively, one may perform separate analyses of the log survival time, conditional on death, and the binary phenotype survival/death. First we set up these phenotypes.

```
y <- listeria$pheno$logSurv
my <- max(y, na.rm=TRUE)
z <- as.numeric(y==my)
y[y==my] <- NA
listeria$pheno$logSurv2 <- y
listeria$pheno$binary <- z
plot(listeria)
```

We use standard interval mapping for the log survival time conditional on death; the results are slightly different from LOD(μ).

```
out.mu <- scanone(listeria, pheno.col=4)
plot(out.mu, out.2p, lodcolumn=c(1,3), chr=c(1,5,13,15), col=c("blue","red"))
```

We can use `scanone` with `model="binary"` to analyze the binary phenotype. Again, the results are only slight different from LOD(p).

```
out.p <- scanone(listeria, pheno.col=5, model="binary")
plot(out.p, out.2p, lodcolumn=c(1,2), chr=c(1,5,13,15), col=c("blue","red"))
```

The argument `pheno.col` in `scanone` can actually take a vector of numeric phenotype values, and not just an indicator to a phenotype column, and so we could have performed the binary trait analysis without first pasting the binary phenotype into the `listeria` object, as follows.

```
out.p.alt <- scanone(listeria, pheno.col=as.numeric(listeria$pheno$T264==264),
                     model="binary")
```

10. A further approach is to use a non-parametric form of interval mapping. R/qtl uses an extension of the Kruskal-Wallis test statistic. Use `scanone` with `model="np"`. In this case, the argument `method` is ignored; the analysis method is much like Haley-Knott regression. If the argument `ties.random=TRUE`, tied phenotypes are ranked at random. If `ties.random=FALSE`, tied phenotypes are given the average rank and a correction is applied to the LOD score.

```
out.np1 <- scanone(listeria, model="np", ties.random=TRUE)
out.np2 <- scanone(listeria, model="np", ties.random=FALSE)

plot(out.np1, out.np2, col=c("blue","red"))
plot(out.2p, out.np1, out.np2, chr=c(1,5,13,15))
```

Note that the significance threshold for the non-parametric genome scan will be quite a bit smaller than that for the two-part model. The two approaches for dealing with ties give basically the same results. Randomizing ties for the non-parametric approach can give quite variable results in the case of a great number of ties, and so we would recommend the use of `ties.random=FALSE` in this case.

Example 4: Covariates in QTL mapping

As a further example, we illustrate the use of covariates in QTL mapping. We consider some simulated backcross data.

1. Get access to the data.

```
data(fake.bc)
summary(fake.bc)
plot(fake.bc)
```

2. Perform genome scans for the two phenotypes without covariates. Here we consider two phenotypes, scanned individually.

```
fake.bc <- calc.genoprob(fake.bc, step=2.5)
out.nocovar <- scanone(fake.bc, pheno.col=1:2)
```

3. Perform genome scans with sex as an additive covariate. Note that the covariates must be numeric. Factors may have to be converted.

```
sex <- fake.bc$pheno$sex
out.acovar <- scanone(fake.bc, pheno.col=1:2, addcovar=sex)
```

Here, the average phenotype is allowed to be different in the two sexes, but the effect of the putative QTL is assumed to be the same in the two sexes.

4. Note that the use of sex as an additive covariate resulted in an increase in the LOD scores for phenotype 1, but resulted in a decreased LOD score at the chr 5 locus for phenotype 2.

```
summary(out.nocovar, threshold=3, format="allpeaks")
summary(out.acovar, threshold=3, format="allpeaks")
plot(out.nocovar, out.acovar, chr=c(2, 5))
plot(out.nocovar, out.acovar, chr=c(2, 5), lodcolumn=2)
```

5. Let us now perform genome scans with sex as an interactive covariate, so that the QTL is allowed to be different in the two sexes.

```
out.icovar <- scanone(fake.bc, pheno.col=1:2, addcovar=sex, intcovar=sex)
```

6. The LOD score in the output is for the comparison of the full model with terms for sex, QTL and QTL×sex interaction to the reduced model with just the sex term. Thus, the degrees of freedom associated with the LOD score is 2 rather than 1, and so larger LOD scores will generally be obtained.

```
summary(out.icovar, threshold=3, format="allpeaks")
plot(out.acovar, out.icovar, chr=c(2,5), col=c("blue", "red"))
plot(out.acovar, out.icovar, chr=c(2,5), lodcolumn=2,
     col=c("blue", "red"))
```

7. The difference between the LOD score with sex as an interactive covariate and the LOD score with sex as an additive covariate concerns the test of the QTL×sex interaction: does the QTL have the same effect in both sexes? The differences, and a plot of the differences, may be obtained as follows.

```
out.sexint <- out.icovar - out.acovar
plot(out.sexint, lodcolumn=1:2, chr=c(2,5), col=c("green", "purple"))
```

The green and purple curves are for the first and second phenotypes, respectively.

8. To test for the QTL×sex interaction, we may perform a permutation test. This is not perfect, as the permutation test eliminates the effect of the QTL, and so we must assume that the distribution of the LOD score for the QTL×sex interaction is the same in the presence of a QTL as under the global null hypothesis of no QTL effect.

The permutation test requires some care. We must perform separate permutations with sex as an additive covariate and with sex as an interactive covariate, but we must ensure, by setting the "seed" for the random number generator, that they use matched permutations of the data.

For the sake of speed, we will use Haley-Knott regression, even though the results above were obtained by standard interval mapping. Also, we will perform just 100 permutations, though 1000 would be preferred.

```
seed <- ceiling(runif(1, 0, 10^8))
set.seed(seed)
operm.acovar <- scanone(fake.bc, pheno.col=1:2, addcovar=sex,
                        method="hk", n.perm=100)
set.seed(seed)
operm.icovar <- scanone(fake.bc, pheno.col=1:2, addcovar=sex,
                        intcovar=sex, method="hk", n.perm=100)
```

Again, the differences concern the QTL×sex interaction.

```
operm.sexint <- operm.icovar - operm.acovar
```

We can use `summary` to get the genome-wide LOD thresholds.

```
summary(operm.sexint, alpha=c(0.05, 0.20))
```

We can also use these results to look at evidence for QTL×sex interaction in our initial scans.

```
summary(out.sexint, perms=operm.sexint, alpha=0.1,
        format="allpeaks", pvalues=TRUE)
```

Example 5: Multiple QTL mapping

We return to the `hyper` data to illustrate some of the more advanced methods for exploring multiple QTL models. Note that the multiple QTL mapping features are currently implemented only for multiple imputation and Haley-Knott regression. We use multiple imputation here, as Haley-Knott regression performs poorly in the case of selective genotyping, which was used for the `hyper` data.

1. First, let us delete everything in our workspace and then re-load the `hyper` data.

```
rm(list=ls())
data(hyper)
```

2. We will be using the multiple imputation method throughout this example, and so we first need to perform the imputations. Recall that more imputations give more precise results, but take more time and memory. To speed things along, we will use only 16 imputations, even though much more would be needed for a definitive analysis. The small number of imputations will make the following results somewhat unpredictable.

```
hyper <- sim.geno(hyper, step=2.5, n.draws=16, err=0.01)
```

3. We first perform a single-QTL genome scan and inspect the results.

```
out1 <- scanone(hyper, method="imp")
plot(out1)
```

As you may recall from the results in Example 1, we have clear evidence for a QTL on chr 4, and strong evidence for a QTL on chr 1. The LOD curve on chr 1 has an interesting double peak, suggestive of possibly two QTL.

There is a hint of further loci on chr 6 and 15 and elsewhere.

4. In the presence of a large-effect QTL, as seen on chr 4, one may wish to repeat the scan, controlling for that locus. This can make the loci with more modest effect more apparent.

A simple (but rough) approach is to pull out the genotypes for a marker near the peak locus, and use that marker as an additive covariate in a single-QTL scan. The peak marker for these data was D4Mit164:

```
max(out1)
```

If the peak LOD score is not at a marker, we may use `find.marker` to identify the marker closest to the LOD peak.

```
find.marker(hyper, 4, 29.5)
```

5. The function `pull.geno` may be used to pull out the genotype data for that marker, but we'll see that most individuals were not typed at D4Mit164.

```
g <- pull.geno(hyper)[,"D4Mit164"]
mean(is.na(g))
```

We may fill in the genotype data using a single imputation, and then use those imputed genotypes as if they were observed. This is not ideal; we'll do this analysis properly below.

```
g <- pull.geno(fill.geno(hyper))[,"D4Mit164"]
```

6. Now we perform the genome scan, controlling for the chr 4 locus. (Note that in an intercross, we would have to re-code the genotype data to be a two-column numeric matrix.)

```
out1.c4 <- scanone(hyper, method="imp", addcovar=g)
```

We can plot the results together with the original genome scan.

```
plot(out1, out1.c4, col=c("blue", "red"))
```

The LOD curve on chr 1 went up quite a bit. (And, of course, the LOD curve on chr 4 went down to near 0.) To see the effect of controlling for the chr 4 locus more clearly, we can plot the differences between the LOD scores.

```
plot(out1.c4 - out1, ylim=c(-3,3))
abline(h=0, lty=2, col="gray")
```

7. We may also look for loci that interact with the chr 4 locus, by including marker D4Mit164 as an interactive covariate.

```
out1.c4i <- scanone(hyper, method="imp", addcovar=g, intcovar=g)
```

The difference between these LOD scores and those obtained with D4Mit164 as a strictly additive covariate indicates evidence for an interaction with the chr 4 locus.

```
plot(out1.c4i - out1.c4)
```

There is nothing particularly interesting here.

8. Now let us perform a 2d scan. This will take a few minutes, as we're doing the scan at a 2.5 cM step size.

```
out2 <- scantwo(hyper, method="imp")
```

9. Let us look at some summaries for the `scantwo` results. Recall that we need to provide five thresholds (see Example 1, item 18 on page 7). We'll ignore the threshold on the epistasis LOD score, T_i, and use the thresholds suggested above.

```
summary(out2, thr=c(6.0, 4.7, Inf, 4.7, 2.6))
```

Your results may be different from mine, since we are using so few imputations, but I see evidence for loci on chr 1 and 4 (which don't appear to interact) and loci on chr 6 and 15 (which do show evidence of epistasis).

This didn't pick up evidence for two QTL on chr 1; we can look directly at the chr 1 results as follows.

```
summary( subset(out2, chr=1) )
```

The LOD score for a second, additive QTL on chr 2 (LOD_{av1}) is ∼1.6; not strong, but not uninteresting.

Evidence for an interaction between loci on chr 7 and 15 had been previously reported. Those results may be inspected as follows.

```
summary( subset(out2, chr=c(7,15)) )
```

Again, this is interesting but not strong.

10. Let us look at some plots of the `scantwo` results. First we make the standard plot with selected chromosomes; the upper triangle contains LOD_i and the lower triangle contains LOD_f.

```
plot(out2, chr=c(1,4,6,7,15))
```

The arguments `lower` and `upper` may be used to change what is plotted in the upper and lower triangles. For example, with `lower="cond-int"`, LOD_{fv1} (evidence for a second QTL, allowing for epistasis) is displayed in the lower triangle, while with `lower="cond-add"`, LOD_{av1} (evidence for a second QTL, assuming no epistasis) is displayed.

```
plot(out2, chr=1, lower="cond-add")
plot(out2, chr=c(6,15), lower="cond-int")
plot(out2, chr=c(7,15), lower="cond-int")
```

Again, evidence for a second QTL on chr 1 is not strong. Evidence for interacting QTL on chr 6 and 15 is quite strong; the 7×15 interaction is not.

11. We can also perform the 2d scan conditional on the chr 4 locus. We'll do this just for chr 1, 6, 7, and 15, to save time.

```
out2.c4 <- scantwo(hyper, method="imp", addcovar=g, chr=c(1,6,7,15))
```

If we look at the same summaries as before, we see decreased evidence for a second QTL on chr 1 and for the 7×15 interaction, but increased evidence for the 6×15 interaction.

```
summary(out2.c4, thr=c(6.0, 4.7, Inf, 4.7, 2.6))
summary( subset(out2.c4, chr=1) )
summary( subset(out2.c4, chr=c(7,15)) )
```

The sort of plots we made before remain interesting.

```
plot(out2.c4)
plot(out2.c4, chr=1, lower="cond-int")
plot(out2.c4, chr=c(6,15), lower="cond-int")
plot(out2.c4, chr=c(7,15), lower="cond-int")
```

We can also look at the differences in the LOD scores, to see how much conditioning on D4Mit164 has affected the results. We need to subset our original results, since we only scanned selected chromosomes in the conditional analysis. The `allow.neg` argument is used to allow negative LOD scores in the `scantwo` plot, as they would generally be replaced with 0.

```
out2sub <- subset(out2, chr=c(1,6,7,15))
plot(out2.c4 - out2sub, allow.neg=TRUE, lower="cond-int")
```

12. Now let us turn to the fit of multiple-QTL models. The function `fitqtl` is used to fit a specific model.

One must first pull out the data on fixed QTL locations using `makeqtl`. We will consider the possibility of two QTL on chr 1, but will ignore the putative QTL on chr 7.

```
qc <- c(1, 1, 4, 6, 15)
qp <- c(43.3, 78.3, 30.0, 62.5, 18.0)
qtl <- makeqtl(hyper, chr=qc, pos=qp)
```

We also create a "formula" which indicates which QTL are to be included in the fit and which interact; the colon (:) indicates an interaction.

```
myformula <- y ~ Q1+Q2+Q3+Q4+Q5 + Q4:Q5
```

We can now fit a model, including the 6×15 interaction, and get a summary of the results.

```
out.fq <- fitqtl(hyper, qtl=qtl, formula = myformula)
summary(out.fq)
```

The first part of the summary describes the overall fit; the LOD score of ∼23 is the $\log_{10}$ likelihood ratio comparing the full model to the null model.

The second part of the summary gives results dropping one term at a time from the model. In the presence of an interaction, if a term included in the interaction is omitted, the interaction is also omitted, and so the rows for the loci on chr 6 and 15 indicate 2 degrees of freedom.

13. One may also use `fitqtl` to get estimated effects of the QTL in the context of the multiple-QTL model. We can use `drop=FALSE`, so that the "drop one at a time" part of the analysis is not performed, and `get.ests=TRUE` to get the estimated effects.

```
out.fq <- fitqtl(hyper, qtl=qtl, formula = myformula, drop=FALSE, get.ests=TRUE)
summary(out.fq)
```

The estimated effects are the differences between the heterozygote and homozygote groups. The interaction effect is the difference between the differences.

14. The function `refineqtl` can be used to refine the estimated positions of the QTL in the context of the multiple-QTL model. A QTL object may be provided, or one may specify the chromosomes and positions, as in `makeqtl`; we'll use the former approach.

```
revqtl <- refineqtl(hyper, qtl=qtl, formula = myformula)
```

The output is a QTL object, like `qtl`; typing its name gives a brief summary.

```
revqtl
```

A couple of the QTL moved, but none by very much.

One may use the `plot.qtl` function to plot the locations of the QTL on the genetic map.

```
plot(revqtl)
```

We can re-run `fitqtl` to get a fit with the new positions; the overall LOD score should have increased slightly. (For me, it increased from 23.0 to 23.7.)

```
out.fq2 <- fitqtl(hyper, qtl=revqtl, formula=myformula)
summary(out.fq2)
```

15. The `scanqtl` function is used to perform general genome scans in the context of a multiple QTL model. It is quite flexible, but not simple to use. For most purposes, one may focus on the functions `addqtl` and `addpair`, which scan for an additional QTL or pair of QTL, respectively, to add to a multiple-QTL model.

We will first use `addqtl` to perform a more precise version of our genome scan conditional on the chr 4 locus. Previously, we had conditioned on imputed genotypes at a marker near the LOD peak on chr 4. With `addqtl` we can do this properly: take proper account of the missing genotype information at the chr 4 locus, rather than taking genotypes from a single imputation as if they had been observed.

The `addqtl` function is much like `fitqtl`, taking a QTL object and formula as arguments. If the formula is omitted, all loci are assumed to be additive. The additional QTL to be scanned may be included in the formula; if there are 5 QTL in the input QTL object, refer to the new QTL as `Q6`. This allows a scan with the new QTL interacting with one or more of the current QTL. If the new QTL is not included in the formula, it is assumed to be strictly additive.

The following performs a scan on all chromosomes, controlling solely for the QTL on chromosome 4. (This is the third QTL in the QTL object `revqtl`, and so we may use as the formula either `y~Q3` or `y~Q3+Q6`. The former is allowed, as an additional additive QTL is assumed.)

```
out1.c4r <- addqtl(hyper, qtl=revqtl, formula=y~Q3)
```

The output is of the same form as produced by the `scanone` function, and so we may use the same plot and summary functions as are used for `scanone` results. (Note that the LOD scores produced by `addqtl` are relative to the model specified in the formula, omitting any terms including the additional QTL being scanned, rather than relative to the null model.).

We may now plot these results with those obtained earlier. The results are actually not too different.

```
plot(out1.c4, out1.c4r, col=c("blue", "red"))
```

It may be more informative to plot the differences

```
plot(out1.c4r - out1.c4, ylim=c(-1.7, 1.7))
abline(h=0, lty=2, col="gray")
```

16. The function `addpair` may be used to perform a 2d scan for an additional pair of QTL, conditioning on the locus on chr 4. If the new QTL are not specified in the formula, a scan as in `scantwo` is performed (that is, for each possible pair of positions for the new QTL, we fit a model in which the two new QTL interact and one in which they are additive).

```
out2.c4r <- addpair(hyper, qtl=revqtl, formula=y~Q3, chr=c(1,6,7,15))
```

The results are of the same form as produced by `scantwo`, and

We can plot the difference between these results and our previous results.

```
plot(out2.c4r - out2.c4, lower="cond-int", allow.neg=TRUE)
```

Again, things have not changed dramatically.

17. The most interesting use of `addqtl` and `addpair` is to scan for additional loci, starting with our five-QTL model (with the loci on 6 and 15 interacting).

First, we scan for an additional additive QTL.

```
out.1more <- addqtl(hyper, qtl=revqtl, formula=myformula)
plot(out.1more)
```

There is not much evidence for an additional QTL.

18. We may next scan for an additional QTL that interacts with one of the QTL in our model, such as the QTL on chr 15. This may be done by indicating the interaction in the formula, using `Q6` to specify the new QTL, since there are five QTL in the `revqtl` object.

```
out.iw4 <- addqtl(hyper, qtl=revqtl, formula=y~Q1+Q2+Q3+Q4+Q5+Q4:Q5+Q6+Q5:Q6)
plot(out.iw4)
```

The LOD scores are just slightly higher, but there are two degrees of freedom in the test. There's nothing particularly exciting here.

19. Now, let us scan for an additional pair. This will take quite a bit of time, so let's focus on a few chromosomes: 2, 5, 7 and 15.

```
out.2more <- addpair(hyper, qtl=revqtl, formula=myformula, chr=c(2,5,7,15))
```

Again, the results are of the form produced by `scantwo`, and so we may use the same plot and summary functions.

```
plot(out.2more, lower="cond-int")
```

Again, there's nothing particularly exciting.

20. Another function of interest is `addint`, for testing the addition of each possible pairwise interactions, one at a time, to a multiple-QTL model.

```
out.ai <- addint(hyper, qtl=revqtl, formula=myformula)
out.ai
```

The results contain one row per interaction, and contain the same sort of information as produced by in the drop-one analysis of `fitqtl`. As the base model (in `myformula`) contains an interaction between the loci on chr 6 and 15, that particular interaction is not tested.

21. We should mention the functions for manipulating QTL objects (produced by `makeqtl`): `addtoqtl`, `dropfromqtl`, `replaceqtl`, and `reorderqtl`.

   If the use of `addqtl` and `addpair` had indicated evidence for additional QTL, one could add them to the QTL object with `addtoqtl`. As input, one provides the cross, the QTL object, and the chromosomes and positions of the QTL to be added.

   ```
   qtl2 <- addtoqtl(hyper, revqtl, 7, 53.6)
   qtl2
   ```

   A QTL may be removed with `dropfromqtl`. One provides either the numeric index within the object, the QTL name, or the chromosome and position of the QTL to be dropped.

   ```
   qtl3 <- dropfromqtl(qtl2, index=2)
   qtl3
   ```

   We can use `replaceqtl` to move a particular QTL to a new position. One must provide the index of the QTL to be replaced.

   ```
   qtl4 <- replaceqtl(hyper, qtl3, index=1, chr=1, pos=50)
   qtl4
   ```

   We use `reorderqtl` to change the order of the loci within a QTL object.

   ```
   qtl5 <- reorderqtl(qtl4, c(1:3,5,4))
   qtl5
   ```

22. Finally, we consider an automated model selection procedure with a stepwise search algorithm, using the function `stepwiseqtl`. The function seeks to optimize a penalized LOD score criterion, which is the LOD score for a model (relative to the null model with no QTL) with penalties on each QTL main effect and a separate penalty on interactions.

   Actually, we include include two penalties on interactions, a light penalty and a heavy penalty. We focus on models with possible pairwise interactions among QTL, and with a hierarchical structure in which the inclusion of an interaction term requires the inclusion of both of the corresponding main effects terms. Such a model may be represented by a graph in which vertices (dots) represent QTL and edges (line segments between the dots) represent interactions between QTL. In the penalized LOD score considered by `stepwiseqtl`, each disconnected component of a model is allowed one light interaction penalty; all other interactions are assigned the heavy penalty.

   The three penalties may be calculated from permutation results with `scantwo`, using the function `calc.penalties`. We will use default penalties derived by computer simulation: (2.69, 2.62, 1.19) for a mouse backcross, or (3.52, 4.28, 2.69) for a mouse intercross. (The penalties are in the order (main, heavy interaction, light interaction).)

   First, let us apply `stepwiseqtl`, considering only additive QTL models (with `additive.only=TRUE`. The algorithm performs forward selection up to a model with a given number of QTL (specified by the argument `max.qtl`; we'll use 6), followed by backward elimination.

   ```
   stepout.a <- stepwiseqtl(hyper, additive.only=TRUE, max.qtl=6)
   stepout.a
   ```

   I obtained a model with two QTL, with one QTL on each of chr 1 and 4.

   Now let's re-run the analysis, allowing for the possibility of interactions among the QTL.

   ```
   stepout.i <- stepwiseqtl(hyper, max.qtl=6)
   stepout.i
   ```

   I obtained a model with four QTL, including one on each of chr 1, 4, 6 and 15, and including an interaction between the loci on chr 6 and 15.

23. Note that all of the above could be performed using Haley-Knott regression rather than multiple imputation. Just three changes need to be made.

   First, one needs to run `calc.genoprob` rather than `sim.geno`, to calculate the QTL genotype probabilities rather than perform imputations.

   Second, in a call to `makeqtl`, use the argument `what="prob"`, so that the genotype probabilities are placed in the object rather than imputations.

   Third, in calls to `fitqtl`, `addqtl`, `addpair`, etc., use `method="hk"`.

Example 6: Internal data structure

Finally, let us briefly describe the rather complicated data structure that R/qtl uses for QTL mapping experiments. This will be rather dull, and will require a good deal of familiarity with the R (or S) language. The choice of data structure required some balance between ease of programming and simplicity for the user interface. The syntax for references to certain pieces of the internal data can become extremely complicated.

1. Get access to some sample data.

   ```
   data(fake.bc)
   ```

2. First, the object has a "class," which indicates that it corresponds to data for an experimental cross, and gives the cross type. By having class `cross`, the functions `plot` and `summary` know to send the data to `plot.cross` and `summary.cross`.

   ```
   class(fake.bc)
   ```

3. Every `cross` object has two components, one containing the genotype data and genetic maps and the other containing the phenotype data.

   ```
   names(fake.bc)
   ```

4. The phenotype data is simply a matrix (more strictly a data.frame) with rows corresponding to individuals and columns corresponding to phenotypes.

   ```
   fake.bc$pheno[1:10,]
   ```

5. The genotype data is a list with components corresponding to chromosomes. Each chromosome has a name and a class. The class for a chromosome is either `"A"` or `"X"`, according to whether it is an autosome or the X chromosome.

   ```
   names(fake.bc$geno)
   sapply(fake.bc$geno, class)
   ```

6. Each component of `geno` contains two components, `data` (containing the marker genotype data) and `map` (containing the positions of the markers, in cM).

   ```
   names(fake.bc$geno[[3]])
   fake.bc$geno[[3]]$data[1:5,]
   fake.bc$geno[[3]]$map
   ```

   That's it for the raw data.

7. When one runs `calc.genoprob`, `sim.geno`, `argmax.geno` or `calc.errorlod`, the output is the input cross object with the derived data attached to each component (the chromosomes) of the `geno` component.

   ```
   names(fake.bc$geno[[3]])
   fake.bc <- calc.genoprob(fake.bc, step=10, err=0.01)
   names(fake.bc$geno[[3]])
   fake.bc <- sim.geno(fake.bc, step=10, n.draws=8, err=0.01)
   names(fake.bc$geno[[3]])
   fake.bc <- argmax.geno(fake.bc, step=10, err=0.01)
   names(fake.bc$geno[[3]])
   fake.bc <- calc.errorlod(fake.bc, err=0.01)
   names(fake.bc$geno[[3]])
   ```

8. Finally, when one runs `est.rf`, a matrix containing the pairwise recombination fractions and LOD scores is added to the cross object.

   ```
   names(fake.bc)
   fake.bc <- est.rf(fake.bc)
   names(fake.bc)
   ```